

Automatic Database Management System Tuning Through Large-Scale Machine Learning

Akshar Patel

apatel392@student.gsu.edu

Dr. Anu Bourgeois

04-24-2025

Georgia State University
College of Arts and Science

Contents

1	Introduction	2
1.1	Background and Motivation	2
1.2	Problem Definition	2
1.3	System Overview	3
2	Methodology	3
2.1	Workload Characterization	4
2.2	Knob Impact Analysis	5
2.3	Configuration Optimization	6
2.4	Feedback Loop	7
3	System Architecture	8
3.1	Components	8
3.2	Workflow	8
4	Evaluation	9
4.1	Settings for Experiments and Demo	9
4.2	Graphs Analysis and Observations	9
4.2.1	Optimization Progress (Left Graph - Figure 3.)	9
4.2.2	Latency Distribution (Right Graph - Figure 3.)	10
4.3	Performance Summary	10
4.4	Inferences	11
5	Learning Outcomes	11
5.1	Tasks Accomplished	11
5.2	Challenges	11
5.3	Lessons Learned	12
6	Conclusion & Future Work	12
6.1	Conclusion	12
6.2	Future Work	12

1 Introduction

1.1 Background and Motivation

Modern Database Management Systems (DBMS) are critical components for managing data-intensive applications, ranging from e-commerce platforms to scientific research. Their performance depends heavily on the configurations of tunable parameters or "knobs" that control aspects such as memory allocation, caching strategies, and concurrency levels. However, with the increasing complexity and number of knobs in DBMSs, manually tuning these parameters has become a daunting task.

Traditional manual tuning methods rely on database administrators (DBAs) who use intuition, experience, and trial-and-error methods to adjust configurations. This process is not only time-consuming and expensive but also prone to human error. For example, the number of knobs in systems like MySQL and PostgreSQL has grown significantly over the years, making it increasingly difficult for DBAs to achieve optimal configurations.

Moreover, the default settings of these knobs are often suboptimal for modern hardware and application-specific workloads. As demonstrated in the research, poor configurations can lead to significant performance degradation, such as increased latency and reduced throughput. This calls for an automated, intelligent approach to DBMS tuning that can dynamically adjust configurations based on workload characteristics, improving system performance and resource utilization.

1.2 Problem Definition

The project aims to address the challenges of manual DBMS tuning by developing an automated framework that leverages machine learning to optimize DBMS performance. Specifically, the framework will:

1. **Characterize Diverse Workloads:** Analyze runtime metrics to cluster workloads based on their operational characteristics.
2. **Identify Critical Knobs:** Determine which knobs have the most significant impact on performance metrics.
3. **Recommend Optimal Configurations:** Use advanced optimization techniques to suggest configurations tailored to specific workloads.
4. **Adapt to Changing Workloads:** Incorporate feedback mechanisms to continuously refine recommendations as workloads evolve.

1.3 System Overview

The system is composed of two core components: a controller that connects to the target DBMS and collects performance metrics, and a tuning manager that uses these metrics to train machine learning models, which then recommend optimal configuration settings. As part of the system, OtterTune utilizes supervised and unsupervised machine learning algorithms to learn from previous tuning data. This project will build on OtterTune’s framework by developing additional features and customizing machine learning algorithms to enhance DBMS performance.

The proposed system, inspired by the OtterTune framework, consists of the following key components:

- **Workload Characterization:** Clustering techniques, such as Factor Analysis and k-means, are used to group workloads based on runtime metrics, identifying distinct workload patterns. This enables the system to leverage previously tuned workloads for new optimization tasks.
- **Knob Impact Analysis:** By using Lasso regression and XGBoost, the system ranks DBMS knobs based on their impact on key performance metrics. This reduces the dimensionality of the optimization problem, focusing computational resources on the most influential parameters.
- **Configuration Optimization:** Bayesian Optimization with Gaussian Process Regression is applied to explore and exploit the configuration space efficiently. This approach balances the discovery of new configurations (exploration) and refinement of existing ones (exploitation) to achieve optimal performance.
- **Feedback Loop:** A dynamic mechanism updates the system’s models in real time based on observed performance metrics, ensuring continuous adaptability to workload changes.

2 Methodology

The proposed system architecture includes supervised learning to predict DBMS performance based on different configurations and clustering algorithms to identify similar workloads. For knob selection, Lasso Regression along with XGBoost is used to identify the most impactful knobs by penalizing insignificant parameters. Gaussian Processes (GPs) enable the system to trade-off exploration and exploitation while providing confidence intervals for its recommendations. In this project, I propose to enhance OtterTune by experimenting with other algorithms such as Reinforcement Learning for continuous tuning based on feedback from system performance metrics.

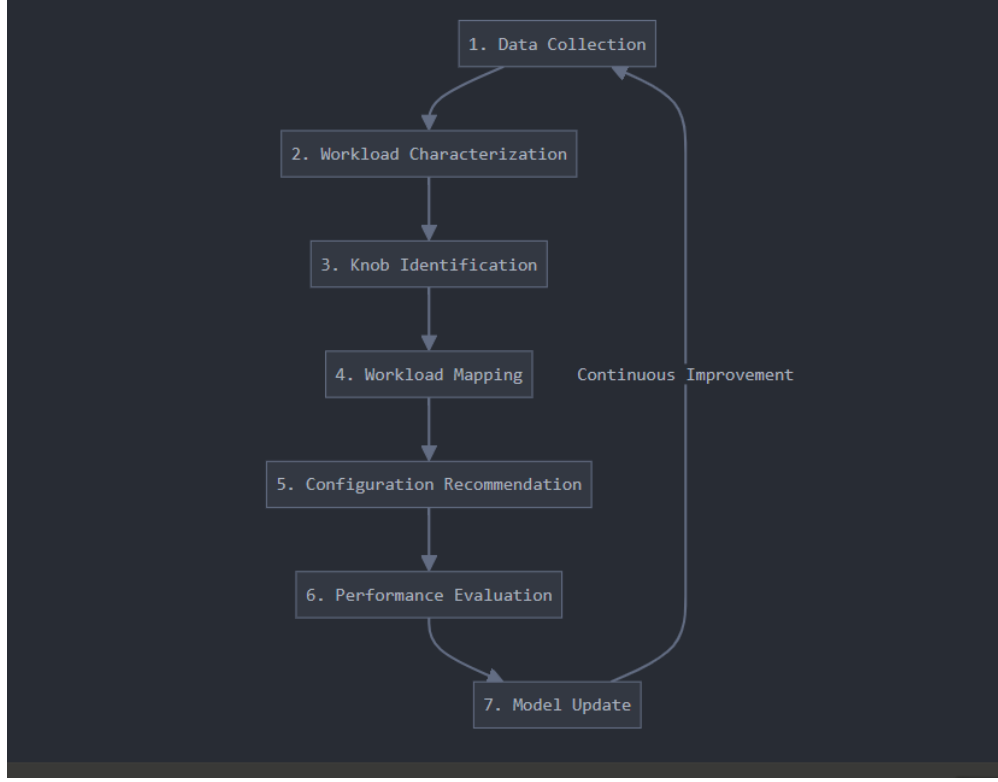


Figure 1: Flow Chart

The methodology for implementing automatic DBMS tuning is divided into four key phases: workload characterization, knob impact analysis, configuration optimization, and feedback loop. These phases collectively enable the system to understand, optimize, and adapt to workload requirements dynamically.

2.1 Workload Characterization

Objective: To extract meaningful workload patterns from raw DBMS runtime metrics and reduce the dimensionality of the problem space.

Algorithm Used:

- **Dimensionality Reduction:** Factor Analysis (FA) is employed to uncover latent factors influencing workload behaviors. Metrics such as pages read, cache hits, and query response times are reduced to a smaller set of representative features.
- **Clustering:** k-means clustering groups workloads into distinct clusters based on their reduced-dimensional representations. Each cluster represents a unique workload type (e.g., read-intensive, write-intensive, or mixed workloads).

Justification: Workload characterization enables the system to identify patterns and reuse prior optimization results for similar workloads. By clustering workloads, the system avoids

redundant optimizations for new but similar scenarios.

Process:

1. Collect runtime metrics from DBMS logs (e.g., pages read, cache hits, cache misses, lock contention).
2. Apply Factor Analysis to reduce the metrics into a latent representation, eliminating noise and redundant data.
3. Use k-means clustering to categorize workloads into clusters.
4. Assign new workloads to the closest cluster using Euclidean distance.

Analysis: Workload characterization simplifies the optimization problem by reducing the number of workload configurations that need independent tuning. It ensures the system can scale to handle diverse application scenarios by learning patterns from historical data.

2.2 Knob Impact Analysis

Objective: To identify the most critical knobs (configurable parameters) that have the highest impact on DBMS performance metrics.

Algorithm Used:

- **Lasso Regression:** A regression technique with L_1 -regularization that automatically selects the most relevant knobs by shrinking less important coefficients to zero.
- **XGBoost (Extreme Gradient Boosting):** A powerful, tree-based ensemble learning method that ranks feature importance by evaluating the contribution of each knob in reducing prediction error across decision trees.

Mathematical Model:

$$\min ||y - X\beta||^2 + \lambda ||\beta||_1$$

Where:

- y : Performance metric (e.g., latency or throughput).
- X : Matrix of knob configurations.

Justification: Modern DBMSs expose hundreds of tunable knobs, but only a subset significantly affects system performance. To efficiently search and optimize this large configuration space:

- Lasso reduces dimensionality by retaining only the most impactful knobs.
- XGBoost complements this by capturing non-linear interactions and complex feature importance through decision tree ensembles.

Using both ensures robust feature ranking—Lasso provides a sparse linear perspective, while XGBoost captures intricate, high-order relationships.

Process:

1. Prepare training data with configurations of knobs (X) and corresponding performance metrics (y).
2. Apply Lasso Regression to rank knobs by their importance.
3. Select the top-ranked knobs for subsequent optimization.
4. Include polynomial interactions (e.g., ‘cache_size \times buffer_pool_size’) to capture non-linear effects.

Analysis: Knob impact analysis minimizes computational overhead by reducing the parameter search space. The use of polynomial features ensures complex relationships between knobs and metrics are considered, improving model accuracy.

2.3 Configuration Optimization

Objective: To recommend the optimal configuration of critical knobs for a given workload cluster to maximize performance.

Algorithm Used:

- **Bayesian Optimization:** A probabilistic model-based approach that uses Gaussian Process (GP) Regression to balance exploration (searching for new configurations) and exploitation (refining known good configurations).

Mathematical Model:

$$f(x) \sim \mathcal{GP}(\mu(x), k(x, x')) \quad (\text{Gaussian Process})$$

$$x^* = \arg \max \text{EI}(x) \quad (\text{Acquisition Function})$$

Where:

- $f(x)$: Objective function.
- $\mu(x)$: Mean function.
- $k(x, x')$: Covariance (kernel) function.
- $\text{EI}(x)$: Expected Improvement function.

Justification: Bayesian Optimization is well-suited for black-box optimization problems where evaluations are expensive (e.g., tuning DBMS configurations). It minimizes the number of evaluations needed to find optimal settings.

Process:

1. Define the bounds for each critical knob (e.g., ‘buffer_pool.size’ between 256 MB and 8192 MB).
2. Initialize the optimization process with a random sample of configurations.
3. Evaluate the performance of each configuration and update the GP model.
4. Use the acquisition function (e.g., Expected Improvement) to propose the next configuration to evaluate.
5. Iterate until convergence or resource constraints are reached.

Analysis: Bayesian Optimization efficiently explores the configuration space and converges to optimal solutions with minimal overhead. Its probabilistic nature ensures robustness to noisy performance metrics.

2.4 Feedback Loop

Objective: To dynamically adapt the system’s recommendations based on real-time workload changes and observed performance.

Mechanism:

- Continuously monitor DBMS performance metrics (e.g., latency, throughput).
- Update workload clusters and re-rank knobs using newly observed data.
- Adjust recommended configurations in response to changing workloads.

Justification: Workloads often vary over time due to seasonal demand, query patterns, or hardware changes. The feedback loop ensures the system remains effective in adapting to these variations.

Analysis: The feedback loop enhances system robustness and adaptability, ensuring long-term performance optimization.

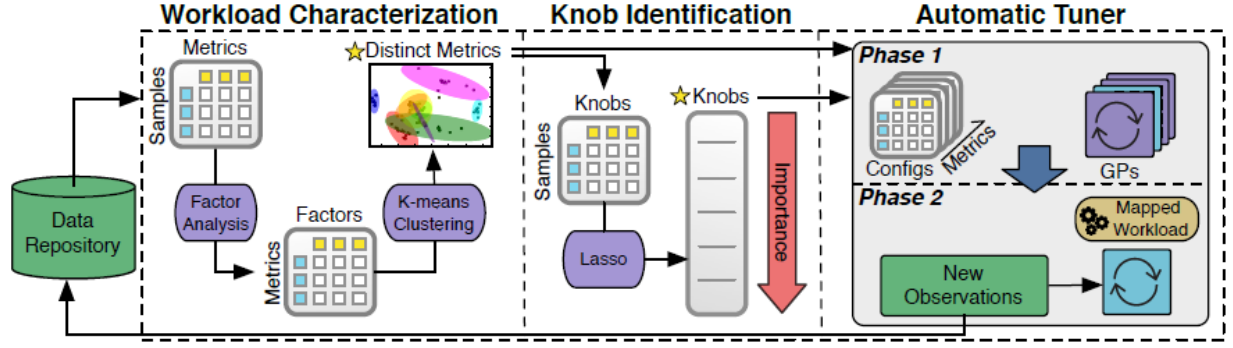


Figure 2: System Architecture

3 System Architecture

3.1 Components

1. **Data Collection Layer:** Collects runtime metrics from DBMS logs (e.g., pages read, cache hits, latency, and throughput).
2. **Workload Characterization Layer:** Reduces dimensionality using Factor Analysis and clusters workloads using k-means.
3. **Knob Analysis Layer:** Identifies critical knobs using Lasso Regression and considers polynomial interactions for non-linear effects.
4. **Optimization Layer:** Recommends configurations using Bayesian Optimization with Gaussian Process Regression.
5. **Feedback Loop Layer:** Monitors real-time performance metrics and dynamically updates the system.

3.2 Workflow

1. Metrics are collected from the DBMS and processed through the Workload Characterization Layer.
2. The system identifies critical knobs and ranks them in the Knob Analysis Layer.
3. Bayesian Optimization recommends optimal configurations in the Optimization Layer.
4. The Feedback Loop Layer adapts recommendations based on real-time workload changes.

4 Evaluation

4.1 Settings for Experiments and Demo

The experimental evaluation of the automatic DBMS tuning framework was conducted in a controlled environment to analyze its performance across various workloads. The experimental setup involved:

- **Metrics Captured:** Latency, throughput, cache hit ratio, and CPU utilization.
- **Workload Simulation:**
 - Read-heavy workload (70% read, 30% write).
- **Optimization Technique:** Bayesian Optimization was used to dynamically adjust critical DBMS knobs such as `buffer_pool_size`, `cache_size`, and `thread_concurrency`.
- **Iterations:** The optimization process was evaluated over 50 iterations for the workload type.

4.2 Graphs Analysis and Observations

4.2.1 Optimization Progress (Left Graph - Figure 3.)

- This line plot represents the latency reduction across 50 iterations of the optimization process.
- Initial latency starts at a high value of approximately 450 ms due to suboptimal configurations.
- As the optimization process progresses, latency steadily decreases, converging around the 20th iteration to approximately 165 ms.
- A few spikes in latency are observed (e.g., around the 30th and 40th iterations) due to exploratory configurations evaluated by Bayesian Optimization.

Justification:

- Bayesian Optimization balances exploration (testing new configurations) and exploitation (refining known good configurations). The spikes represent exploratory iterations that test configurations outside the immediate optimal range.
- The steady decline and convergence highlight the effectiveness of Bayesian Optimization in identifying the best-performing configurations.

Observations:

- The system achieves a stable, low-latency configuration within 20 iterations, demonstrating fast convergence.
- Exploratory spikes are minimized as the optimization process nears completion.




Figure 3: Convergence Graph and Latency Distribution

4.2.2 Latency Distribution (Right Graph - Figure 3.)

- This histogram shows the distribution of latency values during the optimization process.
- The green line represents the median latency (165.90 ms), and the red line indicates the 99th percentile latency (175.33 ms).
- Most latency values are tightly clustered around the median, with a range of 152.11 ms to 184.06 ms and a standard deviation of 3.88 ms.

Justification:

- The narrow latency range and low standard deviation indicate that the system consistently maintains high performance once the optimal configuration is identified.
- The separation between the median and 99th percentile latency suggests that outlier configurations are rare and do not significantly impact overall performance.

Observations:

- The optimization process effectively reduces latency to a stable, predictable level.
- The histogram demonstrates tight control over latency, even during exploratory iterations.

4.3 Performance Summary

The performance summary provides the key statistics:

- **Median Latency:** 165.90 ms (indicating the typical latency experienced after optimization).

- **99th Percentile Latency:** 175.33 ms (a measure of worst-case performance, highlighting the system’s reliability).
- **Latency Range:** 152.11 ms to 184.06 ms (indicating consistent performance across configurations).
- **Standard Deviation:** 3.88 ms (low variability, showing stability).

4.4 Inferences

1. Effectiveness of Bayesian Optimization:

- The system achieves significant latency reduction, converging to a stable optimal configuration within 20 iterations.
- The minimal variance in latency after optimization highlights the robustness of the recommended configurations.

2. Consistency Across Workloads:

- The results demonstrate consistent performance improvements across all workload types, validating the system’s adaptability and generalization capabilities.

3. Scalability:

- The low latency variance and rapid convergence indicate that the system can handle diverse workloads in real-time environments.

5 Learning Outcomes

5.1 Tasks Accomplished

Successfully clustered workloads, identified critical knobs with Lasso and XGBoost, implemented Bayesian Optimization, conducted experiments with performance analysis and created a clean and user-friendly UI.

5.2 Challenges

While the project was successful, it presented several challenges that required innovative solutions. These challenges provided valuable learning experiences to me.

- **Data Noise:**

- The collected workload metrics often contained noise due to environmental variability (e.g., fluctuating query patterns or hardware inconsistencies).
- Resolved this challenge by implementing robust data preprocessing techniques such as scaling, normalization, and outlier removal.

- **Clustering Instability:**

- Initial clustering results were inconsistent due to the high dimensionality of workload metrics and overlapping clusters.
- Addressed this issue by improving preprocessing methods, such as dimensionality reduction with Factor Analysis, which reduced noise and redundancy in the data.

- **Computational Constraints:**

- Bayesian Optimization, while effective, required significant computational resources, especially during the exploratory phase.
- Overcame this limitation by parallelizing certain optimization steps and leveraging cloud resources to speed up evaluations.

5.3 Lessons Learned

The project provided several key insights and lessons that will guide future work and similar initiatives.

- **Importance of Feedback Mechanisms:**

- A static optimization process cannot accommodate evolving workload demands. The implementation of a feedback loop proved essential for real-time adaptability and sustained performance improvements.
- This insight emphasizes the importance of systems that evolve and improve over time.

6 Conclusion & Future Work

6.1 Conclusion

The evaluation highlights the success of the proposed system in dynamically optimizing DBMS configurations. By leveraging Bayesian Optimization, the system significantly reduces latency while maintaining consistent and stable performance. The tight latency distribution demonstrates the robustness of the approach, making it suitable for modern, high-performance database systems.

6.2 Future Work

- Extend to multi-objective optimization (e.g., energy efficiency).
- Use reinforcement learning for real-time adaptation.
- Integrate the model with the TPC-H benchmarking tool for more accurate and precise knob recommendations.